

# Clustering Algorithms and High Dimensional Data

Max Mihailescu

December 4th, 2019

# Outline

Introduction

Two Basic Approaches

The Curse of Dimensionality

The Optigrid algorithm

Outlook


# Introduction

Given a set of data-points, the goal of clustering is to group similar points together and have a high dissimilarity between the groups.

A data-point is characterized by a feature vector  $X = \{x_1, x_2, \dots, x_d\}$ , where  $d$  is its dimension.

Various notions of similarity:

- ▶ Distance between objects
- ▶ Density around a given point
- ▶ Distribution of features



---

<sup>1</sup><https://towardsdatascience.com/introduction-to-image-segmentation-with-k-means-clustering-83fd0a9e2fc3>

# K-Means

**Goal:** Partition the data into  $k$  sets, so as to minimize the squared distances between points within a cluster.

# K-Means


**Parameters:** *num\_clusters*

- 1 Initialize *num\_clusters* cluster centers
- 2 assign points to cluster based on nearest center
- 3 **while** *not centres converged* **do**
- 4     Recalculate centers
- 5     Reassign points to new cluster
- 6 **end**

# K-Means

**Parameters:** *num\_clusters*


- 1 Initialize *num\_clusters* cluster centers
- 2 assign points to cluster based on nearest center
- 3 **while** *not centres converged* **do**
- 4 |   Recalculate centers
- 5 |   Reassign points to new cluster
- 6 **end**



# K-Means

**Parameters:** *num\_clusters*


- 1 Initialize *num\_clusters* cluster centers
- 2 assign points to cluster based on nearest center
- 3 **while** *not centres converged* **do**
- 4     Recalculate centers
- 5     Reassign points to new cluster
- 6 **end**



# K-Means

**Parameters:** *num\_clusters*

- 1 Initialize *num\_clusters* cluster centers
- 2 assign points to cluster based on nearest center
- 3 **while** *not centres converged* **do**
- 4     Recalculate centers
- 5     Reassign points to new cluster
- 6 **end**






# K-Means

**Parameters:** *num\_clusters*


- 1 Initialize *num\_clusters* cluster centers
- 2 assign points to cluster based on nearest center
- 3 **while** *not centres converged* **do**
- 4     Recalculate centers
- 5     Reassign points to new cluster
- 6 **end**



# K-Means

**Parameters:** *num\_clusters*


- 1 Initialize *num\_clusters* cluster centers
- 2 assign points to cluster based on nearest center
- 3 **while** *not centres converged* **do**
  - 4 | Recalculate centers
  - 5 | Reassign points to new cluster
- 6 **end**



# K-Means

**Parameters:** *num\_clusters*


- 1 Initialize *num\_clusters* cluster centers
- 2 assign points to cluster based on nearest center
- 3 **while** *not centres converged* **do**
  - 4 | Recalculate centers
  - 5 | Reassign points to new cluster
- 6 **end**



# K-Means

**Parameters:** *num\_clusters*


- 1 Initialize *num\_clusters* cluster centers
- 2 assign points to cluster based on nearest center
- 3 **while** *not centres converged* **do**
- 4     Recalculate centers
- 5     Reassign points to new cluster
- 6 **end**



# K-Means

**Parameters:** *num\_clusters*


- 1 Initialize *num\_clusters* cluster centers
- 2 assign points to cluster based on nearest center
- 3 **while** *not centres converged* **do**
- 4     Recalculate centers
- 5     Reassign points to new cluster
- 6 **end**



# K-Means

**Parameters:** *num\_clusters*

- 1 Initialize *num\_clusters* cluster centers
- 2 assign points to cluster based on nearest center
- 3 **while** *not centres converged* **do**
- 4 |   Recalculate centers
- 5 |   Reassign points to new cluster
- 6 **end**



# K-Means

- ▶ Number of clusters is an input parameter
- ▶ Based on a notion of distance
- ▶ Tends to find spherical clusters of the same size
- ▶ Unstable with respect to initial starting points

# DBScan

**Goal:** Find connected regions in space that are highly populated.

---


<sup>a</sup>Chire, CC BY-SA 3.0,  
<https://commons.wikimedia.org/w/index.php?curid=17045963>



# DBScan

**Goal:** Find connected regions in space that are highly populated.


- ▶ **Core points:** Have at least  $minPts$  neighbors within distance  $eps$
- ▶ **Directly reachable points:** At least one core point within distance  $eps$



# DBScan

**Goal:** Find connected regions in space that are highly populated.

- ▶ **Core points:** Have at least  $minPts$  neighbors within distance  $eps$
- ▶ **Directly reachable points:** At least one core point within distance  $eps$




$B$  is reachable from  $A$ , if there exists a path of following directly reachable points.

# DBScan

**Goal:** Find connected regions in space that are highly populated.

- **Core points:** Have at least  $minPts$  neighbors within distance  $eps$
- **Directly reachable points:** At least one core point within distance  $eps$



$B$  is reachable from  $A$ , if there exists a path of following directly reachable points.

$B$  and  $C$  are density connected if, they are reachable from a common point  $A$ .


---

<sup>a</sup>Chire, CC BY-SA 3.0,  
<https://commons.wikimedia.org/w/index.php?curid=17045963>

# DBScan

**Parameters:**  $minPts$ ,  $eps$


- 1 Find all neighbors closer than  $eps$  of every point
- 2 Identify core points with more than  $minPts$  neighbors
- 3 Create clusters of neighboring core points
- 4 Assign the remaining points to a nearby cluster or noise



# DBScan

**Parameters:**  $minPts$ ,  $eps$


- 1 Find all neighbors closer than  $eps$  of every point
- 2 Identify core points with more than  $minPts$  neighbors
- 3 Create clusters of neighboring core points
- 4 Assign the remaining points to a nearby cluster or noise



# DBScan

**Parameters:**  $minPts$ ,  $eps$

- 1 Find all neighbors closer than  $eps$  of every point
- 2 Identify core points with more than  $minPts$  neighbors
- 3 Create clusters of neighboring core points
- 4 Assign the remaining points to a nearby cluster or noise



# DBScan

- ▶ Groups by point density in a neighborhood that is controlled using input parameters
- ▶ Finds arbitrarily shaped clusters
- ▶ Has a notion of noise
- ▶ Generally stable between runs
- ▶ Parameters require knowledge of the data

# The Curse of Dimensionality

Data quickly becomes sparse in high dimensions:

- ▶ Consider a problem with 50 dimensions of binary features and  $10^{12}$  observations.
- ▶ Then we still only have examples for  $\frac{10^{12}}{2^{50}} \approx 0.089\%$  of categories.



# The Curse of Dimensionality

Data quickly becomes sparse in high dimensions:


- ▶ Consider a problem with 50 dimensions of binary features and  $10^{12}$  observations.
- ▶ Then we still only have examples for  $\frac{10^{12}}{2^{50}} \approx 0.089\%$  of categories.

**Problem 1:** It becomes impossible to correctly estimate distributions.

# The Curse of Dimensionality

Our intuition fails in higher dimensions:


- ▶ Consider uniformly distributed data on a cube in  $d$  dimensions with side length 1.
- ▶ How big can we make  $a$  so that more than 95% of the probability mass lies inside the dotted region?



# The Curse of Dimensionality

Our intuition fails in higher dimensions:

- ▶ Consider uniformly distributed data on a cube in  $d$  dimensions with side length 1.
- ▶ How big can we make  $a$  so that more than 95% of the probability mass lies inside the dotted region?




$$0.95 \leq P(X \in \text{dotted})$$

# The Curse of Dimensionality

Our intuition fails in higher dimensions:

- ▶ Consider uniformly distributed data on a cube in  $d$  dimensions with side length 1.
- ▶ How big can we make  $a$  so that more than 95% of the probability mass lies inside the dotted region?




$$0.95 \leq P(X \in \text{dotted}) = P(X \in \text{bigcube}) - P(X \in \text{smallcube}) = 1 - a^d$$

# The Curse of Dimensionality

Our intuition fails in higher dimensions:

- ▶ Consider uniformly distributed data on a cube in  $d$  dimensions with side length 1.
- ▶ How big can we make  $a$  so that more than 95% of the probability mass lies inside the dotted region?




$$0.95 \leq P(X \in \text{dotted}) = P(X \in \text{bigcube}) - P(X \in \text{smallcube}) = 1 - a^d$$
$$\Leftrightarrow a \leq \sqrt[d]{0.05}$$

# The Curse of Dimensionality


Our intuition fails in higher dimensions:

- ▶ Consider uniformly distributed data on a cube in  $d$  dimensions with side length 1.
- ▶ How big can we make  $a$  so that more than 95% of the probability mass lies inside the dotted region?



$$0.95 \leq P(X \in \text{dotted}) = P(X \in \text{bigcube}) - P(X \in \text{smallcube}) = 1 - a^d$$


$$\Leftrightarrow a \leq \sqrt[d]{0.05}$$



# The Curse of Dimensionality

Our intuition fails in higher dimensions:


- ▶ Consider uniformly distributed data on a cube in  $d$  dimensions with side length 1.
- ▶ How big can we make  $a$  so that more than 95% of the probability mass lies inside the dotted region?



$$0.95 \leq P(X \in \text{dotted}) = P(X \in \text{bigcube}) - P(X \in \text{smallcube}) = 1 - a^d$$

$$\Leftrightarrow a \leq \sqrt[d]{0.05}$$

**Problem 2:** Regions with a high density can be empty.



# Optigrid

A grid based approach to overcome the problems of distance and density in high dimensional spaces.

Hinneburg and Keim, *Optimal Grid-Clustering: Towards Breaking the Curse of Dimensionality in High-Dimensional Clustering*

**Goal:** Partition the space into grid-cells such that dense regions do not get cut in half.




# Optigrid

A grid based approach to overcome the problems of distance and density in high dimensional spaces.

Hinneburg and Keim, *Optimal Grid-Clustering: Towards Breaking the Curse of Dimensionality in High-Dimensional Clustering*

**Goal:** Partition the space into grid-cells such that dense regions do not get cut in half.

Key Observation: If the density at a point in one coordinate is low, it is low at all points in the high dimensional space that get projected onto that point.



# Optigrid

**Input:** Data

**Parameters:**  $q$ ,  $maxCutScore$ ,  $noiseLevel$


```
1 Calculate Projections to lower dimensional space
2 Approx. Densities and find peaks above  $noiseLevel$ 
3 Find  $q$  best cutting planes with score  $\leq maxCutScore$ 
4 if cutting planes were found then
5   | Divide data into subgrids  $C_1, \dots, C_n$ 
6   | foreach subgrid do
7   |   | Optigrid(subgrid)
8   | end
9 else
10  | Mark subgrid as cluster
11 end
```

# Optigrid

**Input:** Data

**Parameters:**  $q$ ,  $maxCutScore$ ,  $noiseLevel$

- 1 Calculate Projections to lower dimensional space
- 2 Approx. Densities and find peaks above  $noiseLevel$
- 3 Find  $q$  best cutting planes with score  $\leq maxCutScore$
- 4 **if** *cutting planes were found* **then**
  - 5 | Divide data into subgrids  $C_1, \dots, C_n$
  - 6 | **foreach** *subgrid* **do**
  - 7 | | Optigrid(subgrid)
  - 8 | **end**
- 9 **else**
- 10 | Mark subgrid as cluster
- 11 **end**




# Optigrid

**Input:** Data

**Parameters:**  $q$ ,  $maxCutScore$ ,  $noiseLevel$

- 1 Calculate Projections to lower dimensional space
- 2 Approx. Densities and find peaks above  $noiseLevel$
- 3 Find  $q$  best cutting planes with score  $\leq maxCutScore$
- 4 **if** *cutting planes were found* **then**
  - 5 | Divide data into subgrids  $C_1, \dots, C_n$
  - 6 | **foreach** *subgrid* **do**
  - 7 | | Optigrid(subgrid)
  - 8 | **end**
- 9 **else**
- 10 | Mark subgrid as cluster
- 11 **end**




# Optigrid

**Input:** Data

**Parameters:**  $q$ ,  $maxCutScore$ ,  $noiseLevel$

- 1 Calculate Projections to lower dimensional space
- 2 Approx. Densities and find peaks above  $noiseLevel$
- 3 Find  $q$  best cutting planes with score  $\leq maxCutScore$
- 4 **if** *cutting planes were found* **then**
  - 5 | Divide data into subgrids  $C_1, \dots, C_n$
  - 6 | **foreach** *subgrid* **do**
  - 7 | | Optigrid(subgrid)
  - 8 | **end**
- 9 **else**
- 10 | Mark subgrid as cluster
- 11 **end**




# Optigrid

**Input:** Data

**Parameters:**  $q$ ,  $maxCutScore$ ,  $noiseLevel$

- 1 Calculate Projections to lower dimensional space
- 2 Approx. Densities and find peaks above  $noiseLevel$
- 3 Find  $q$  best cutting planes with score  $\leq maxCutScore$
- 4 **if** *cutting planes were found* **then**
  - 5 | Divide data into subgrids  $C_1, \dots, C_n$
  - 6 | **foreach** *subgrid* **do**
  - 7 | | Optigrid(subgrid)
  - 8 | **end**
- 9 **else**
- 10 | Mark subgrid as cluster
- 11 **end**




# Optigrid

**Input:** Data

**Parameters:**  $q$ ,  $maxCutScore$ ,  $noiseLevel$

- 1 Calculate Projections to lower dimensional space
- 2 Approx. Densities and find peaks above  $noiseLevel$
- 3 Find  $q$  best cutting planes with score  $\leq maxCutScore$
- 4 **if** *cutting planes were found* **then**
  - 5 | Divide data into subgrids  $C_1, \dots, C_n$
  - 6 | **foreach** *subgrid* **do**
  - 7 | | Optigrid(subgrid)
  - 8 | **end**
- 9 **else**
- 10 | Mark subgrid as cluster
- 11 **end**




# Optigrid

**Input:** Data

**Parameters:**  $q$ ,  $maxCutScore$ ,  $noiseLevel$

- 1 Calculate Projections to lower dimensional space
- 2 Approx. Densities and find peaks above  $noiseLevel$
- 3 Find  $q$  best cutting planes with score  $\leq maxCutScore$
- 4 **if** *cutting planes were found* **then**
  - 5 | Divide data into subgrids  $C_1, \dots, C_n$
  - 6 | **foreach** *subgrid* **do**
  - 7 | | Optigrid(subgrid)
  - 8 | **end**
- 9 **else**
- 10 | Mark subgrid as cluster
- 11 **end**






# Optigrid

**Input:** Data

**Parameters:**  $q$ ,  $maxCutScore$ ,  $noiseLevel$

- 1 Calculate Projections to lower dimensional space
- 2 Approx. Densities and find peaks above  $noiseLevel$
- 3 Find  $q$  best cutting planes with score  $\leq maxCutScore$
- 4 **if** *cutting planes were found* **then**
  - 5 | Divide data into subgrids  $C_1, \dots, C_n$
  - 6 | **foreach** *subgrid* **do**
  - 7 | | Optigrid(subgrid)
  - 8 | **end**
- 9 **else**
- 10 | Mark subgrid as cluster
- 11 **end**




# Optigrid

**Input:** Data

**Parameters:**  $q$ ,  $maxCutScore$ ,  $noiseLevel$

- 1 Calculate Projections to lower dimensional space
- 2 Approx. Densities and find peaks above  $noiseLevel$
- 3 Find  $q$  best cutting planes with score  $\leq maxCutScore$
- 4 **if** *cutting planes were found* **then**
  - 5 | Divide data into subgrids  $C_1, \dots, C_n$
  - 6 | **foreach** *subgrid* **do**
  - 7 | | Optigrid(subgrid)
  - 8 | **end**
- 9 **else**
- 10 | Mark subgrid as cluster
- 11 **end**




# Optigrid

**Input:** Data

**Parameters:**  $q$ ,  $maxCutScore$ ,  $noiseLevel$

- 1 Calculate Projections to lower dimensional space
- 2 Approx. Densities and find peaks above  $noiseLevel$
- 3 Find  $q$  best cutting planes with score  $\leq maxCutScore$
- 4 **if** *cutting planes were found* **then**
  - 5 | Divide data into subgrids  $C_1, \dots, C_n$
  - 6 | **foreach** *subgrid* **do**
  - 7 | | Optigrid(subgrid)
  - 8 | **end**
- 9 **else**
- 10 | Mark subgrid as cluster
- 11 **end**




# Optigrid

**Input:** Data

**Parameters:**  $q$ ,  $maxCutScore$ ,  $noiseLevel$

- 1 Calculate Projections to lower dimensional space
- 2 Approx. Densities and find peaks above  $noiseLevel$
- 3 Find  $q$  best cutting planes with score  $\leq maxCutScore$
- 4 **if** *cutting planes were found* **then**
  - 5 | Divide data into subgrids  $C_1, \dots, C_n$
  - 6 | **foreach** *subgrid* **do**
  - 7 | | Optigrid(subgrid)
  - 8 | **end**
- 9 **else**
- 10 | Mark subgrid as cluster
- 11 **end**



# Optigrid

- ▶ Divides the input space into an irregular grid
- ▶ Calculations happen in low dimensions
- ▶ Finds clusters that can be enclosed in a rectangle
- ▶ Parameters can be chosen by analyzing the density distributions

# Outlook

Can we find any clusters in the parameter settings of the Linac3 ion source?

